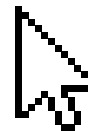




Workshop: Intro to SQL

DSI SUDS Skills Day



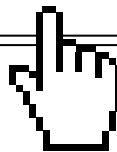
Ciara Zogheib



Hello!

- Ciara Zogheib (She/Her)
- Starting my PhD in information this fall
- Working as a data scientist with government for ~2 years
- First learned SQL in a tiny island village in 2 weeks





01 Background Info

What's an RDBMS, what's SQL, SQL 'dialects'

02 Best Practices

Things to keep in mind when we're writing queries

03 Playing with SQL

Basic queries, exploring a sample database



Today's Goal:

Be able to go into a relational database and pull out specific subsets of data using SQL.



01

Background Info





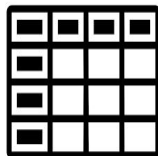
Data

- We can think of data as symbols representing observed properties of objects, events, or the world around us
- Data can be numbers, words, stories, colours, sounds - any systematically collected and organized observation
- **Unstructured data** = documents or images, individual files
- **Structured data** = data in tables, Excel spreadsheet

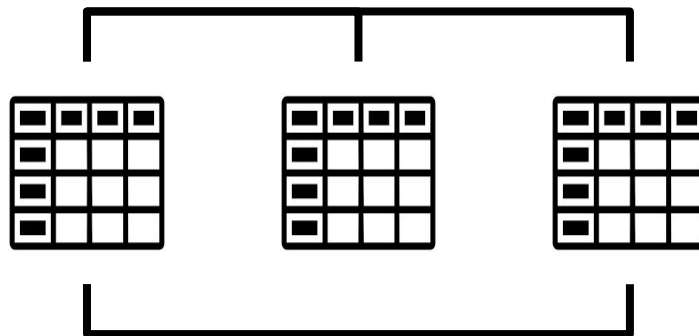
Today, we focus on structured tabular data



Tables VS Databases



■	■	■	■
■			
■			
■			



- A **relational database management system (RDBMS)** lets us store, view, and query multiple tables and the relationships between them
- There are many different types of RDBMS (e.g. MySQL, Oracle, SQLite, MS Access, etc.)



Related Tables

- Related tables are connected by primary and foreign keys
- **Primary key** = an attribute that uniquely identifies a row/record (e.g. student id number or an automatically generated incremental integer number)
- **Foreign key** = when a primary key is referenced in another table (e.g. the table of teachers' students would reference student id numbers)
- Primary and foreign keys help us **normalize** our databases



SQL

- SQL (**Structured Query Language**) lets us ‘query’ (access, extract from, and manipulate) the data in our RDBMS
- SQL can help us extract and combine data from multiple related tables at once, or with conditions applied
- Different RDBMSs sometimes use different SQL syntax, but a general understanding of SQL is easily transferable (different dialects, not different languages)



Today's Most Controversial Slide

How should you
pronounce SQL?



...however you
want.*

*within reason



SQL queries can be simple...

```
SELECT * FROM studentstable WHERE grade = 2;
```



...or complex

```
WITH tmp_1 AS
(
  SELECT Calc1 =
    ( (SELECT TOP 1 DataValue
      FROM (
        SELECT TOP 50 PERCENT DataValue
        FROM SOMEDATA
        WHERE DataValue IS NOT NULL
        ORDER BY DataValue
      ) AS A
    ORDER BY DataValue DESC
    ) +

    (SELECT TOP 1 DataValue
    FROM (
      SELECT TOP 50 PERCENT DataValue
```



...and integrated with other code

```
def create_server_connection(host_name, user_name, user_password):  
    connection = None  
    try:  
        connection = mysql.connector.connect(  
            host=host_name,  
            user=user_name,  
            passwd=user_password  
        )  
        print("MySQL Database connection successful")  
    except Error as err:  
        print(f"Error: '{err}'")  
  
    return connection
```



02

Best Practices





Best Practices

- Write SQL keywords uppercase and table/column names lowercase
 - `SELECT column FROM table;` **not** `select column from table;`
- Indent and use spaces to make long SQL queries readable
- Explicitly specify what columns you want to select
 - `SELECT name, age FROM table;` **not** `SELECT * FROM table`
- Consider using multiple queries instead of one monster query
 - When working with large quantities of data, multiple queries can slow processing and decrease performance - use judgment!
- Comment your code!



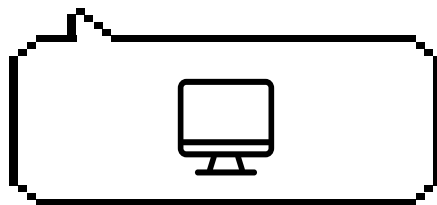
Best Practices

- Have a backup of your data if you are adding, removing, or changing data using SQL queries!
 - A copy of the table in the same database and/or
 - A copy of the table in a different backup database and/or
 - A backup of the entire database
- Double check your data after executing some change



03

Playing with SQL





Oracle SQL Developer : demodb

File Edit View Navigate Run Versioning Tools Help

Connections x Reports x Start Page x demodb x

Connections

- demodb
 - Tables (Filtered)
 - Views
 - Editing Views
 - Indexes
 - Packages
 - Procedures
 - Functions
 - Queues
 - Queues Tables
 - Triggers
 - Crossed Join Triggers
 - Types
 - Sequences
 - Materialized Views
 - Materialized Views Logs
 - Synonyms
 - Public Synonyms
 - Database Links
 - Public Database Links
 - Directories
 - Editions
 - Application Express
 - Java
 - XML Schemas
 - XML DB Repository
 - Scheduler
 - Recycle Bin
 - Other Users
 - Cloud Connections

Worksheet Query Builder

```
select * from v
SYS.v_$active_services
SYS.v_$active_sess_pool_mth
SYS.v_$active_session_history
SYS.v_$advisor_progress
SYS.v_$alert_types
SYS.v_$aq
SYS.v_$aq1
SYS.v_$archive
SYS.v_$archive_dest
SYS.v_$archive_dest_status
SYS.v_$archive_gap
SYS.v_$archive_processes
SYS.v_$archived_log
SYS.v_$ash_info
SYS.v_$asm_acfsanaphots
SYS.v_$asm_acfsavolumes
SYS.v_$asm_alias
SYS.v_$asm_attribute
SYS.v_$asm_client
SYS.v_$asm_disk
SYS.v_$asm_disk_ioatst
SYS.v_$asm_disk_stat
SYS.v_$asm_diskgroup
SYS.v_$asm_diskgroup_stat
SYS.v_$asm_file
SYS.v_$asm_filesystem
SYS.v_$asm_operation
SYS.v_$asm_template
SYS.v_$asm_user
SYS.v_$asm_usergroup
SYS.v_$asm_usergroup_member
SYS.v_$asm_volume
SYS.v_$asm_volume_stat
SYS.v_$aw_aggregate_op
SYS.v_$aw_allocate_op
SYS.v_$aw_calc
SYS.v_$aw_longops
SYS.v_$aw_olap
SYS.v_$aw_session_info
SYS.v_$backup
SYS.v_$backup_archivelog_details
```

BLOG_COMPRESSION SWITCHOVER_STATUS DATAGUARD_BROKER GUARD_STATUS SUF

NOT ALLOWED	DISABLED	NONE	NO
-------------	----------	------	----

Line 1 Column 16 | Insert | Modified | Windows: CR/LF Editing



Schema_Music.mwb - MySQL Workbench

Local instance 3306 x MySQL Model x EER Diagram x

Management Schemas MySQL Schema_Music_Example x

Limit to 1000 rows

SCHEMAS

Filter objects

Music

- Tables
 - Albums
 - Columns
 - Indexes
 - Foreign Keys
 - Triggers
 - Artists
 - Columns
 - Indexes
 - Foreign Keys
 - Triggers
 - Genres
 - Columns
 - Indexes
 - Foreign Keys
 - Triggers
- Views
- Stored Procedures
- Functions
- sys

Object Info Session

Schema: Music

```
1 -- MySQL Script generated by MySQL Workbench
2 -- Mon May 30 11:25:32 2016
3 -- Model: New Model Version: 1.0
4 -- MySQL Workbench Forward Engineering
5
6 • SET @OLD_UNIQUE_CHECKS=@@UNIQUE_CHECKS, UNIQUE_CHECKS=0;
7 • SET @OLD_FOREIGN_KEY_CHECKS=@@FOREIGN_KEY_CHECKS, FOREIGN_KEY_CHECKS=0;
8 • SET @OLD_SQL_MODE=@@SQL_MODE, SQL_MODE='TRADITIONAL,ALLOW_INVALID_DATES';
9
10 -----
11 -- Schema Music
12 -----
13 • DROP SCHEMA IF EXISTS 'Music' ;
14
15 -----
16 -- Schema Music
17 -----
18 • CREATE SCHEMA IF NOT EXISTS 'Music' DEFAULT CHARACTER SET utf8 ;
19 • USE 'Music' ;
20
```

100% 45.1

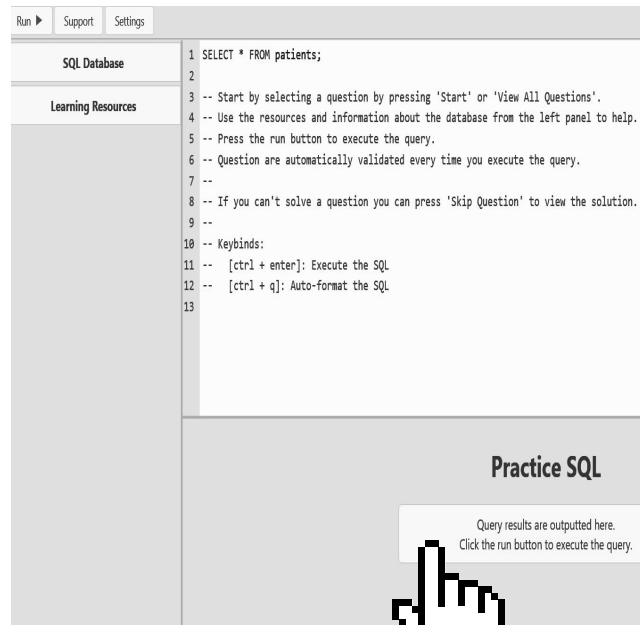
Action	Time	Action	Response	Duration / Fetch Time
4	11:27:13	DROP SCHEMA IF EXISTS 'Music'	0 row(s) affected, 1 warning(s): 1008 Can't drop data...	0.00034 sec
5	11:27:13	CREATE SCHEMA IF NOT EXISTS '...	1 row(s) affected	0.0012 sec
6	11:27:13	USE 'Music'	0 row(s) affected	0.00030 sec
7	11:27:13	DROP TABLE IF EXISTS 'Music'.Artists	0 row(s) affected, 1 warning(s): 1061 Unknown table...	0.00032 sec
8	11:27:13	CREATE TABLE IF NOT EXISTS 'Mu...	0 row(s) affected	0.027 sec
9	11:27:13	DROP TABLE IF EXISTS 'Music'.Ge...	0 row(s) affected, 1 warning(s): 1061 Unknown table...	0.00037 sec
10	11:27:13	CREATE TABLE IF NOT EXISTS 'Mu...	0 row(s) affected	0.026 sec
11	11:27:14	DROP TABLE IF EXISTS 'Music'.Alb...	0 row(s) affected, 1 warning(s): 1061 Unknown table...	0.00029 sec
12	11:27:14	CREATE TABLE IF NOT EXISTS 'Mu...	0 row(s) affected	0.019 sec
13	11:27:14	SET SQL_MODE=@OLD_SQL_MODE	0 row(s) affected	0.00025 sec
14	11:27:14	SET FOREIGN_KEY_CHECKS=@OLD...	0 row(s) affected	0.00026 sec
15	11:27:14	SET UNIQUE_CHECKS=@OLD_UNI...	0 row(s) affected	0.00028 sec

Query Completed



www.sql-practice.com

We're going to use a generic
online sample RDBMS to practice
basic SQL queries



The screenshot shows the SQL Practice web application interface. At the top, there are three tabs: 'Run', 'Support', and 'Settings'. Below these, the interface is split into two main sections. The left section contains a sidebar with 'SQL Database' and 'Learning Resources' links. The right section is a large text area for writing SQL queries. It contains a sample query: `1 SELECT * FROM patients;` followed by several lines of comments explaining the interface and providing keyboard shortcuts. At the bottom right of the interface, there is a 'Practice SQL' button and a text box for the results, which currently says 'Query results are outputted here. Click the run button to execute the query.'

```
1 SELECT * FROM patients;
2
3 -- Start by selecting a question by pressing 'Start' or 'View All Questions'.
4 -- Use the resources and information about the database from the left panel to help.
5 -- Press the run button to execute the query.
6 -- Question are automatically validated every time you execute the query.
7 --
8 -- If you can't solve a question you can press 'Skip Question' to view the solution.
9 --
10 -- Keybinds:
11 -- [ctrl + enter]: Execute the SQL
12 -- [ctrl + q]: Auto-format the SQL
13
```

Practice SQL

Query results are outputted here.
Click the run button to execute the query.





Operators

Operator	Description
----------	-------------

=	Equal to
---	----------

>	Greater than
---	--------------

<	Less than
---	-----------

>=	Greater than or equal to
----	--------------------------

<=	Less than or equal to
----	-----------------------

<>	Not equal to
----	--------------

Operator	Description
----------	-------------

ALL	TRUE if all of the subquery values meet the condition
-----	---

AND	TRUE if all the conditions separated by AND is TRUE
-----	---

ANY	TRUE if any of the subquery values meet the condition
-----	---

BETWEEN	TRUE if the operand is within the range of comparisons
---------	--

EXISTS	TRUE if the subquery returns one or more records
--------	--

IN	TRUE if the operand is equal to one of a list of expressions
----	---

LIKE	TRUE if the operand matches a pattern
------	---------------------------------------

NOT	Displays a record if the condition(s) is NOT TRUE
-----	---

OR	TRUE if any of the conditions separated by OR is TRUE
----	---

SOME	TRUE if any of the subquery values meet the condition
------	---



SQL Functions

- `COUNT()` gives the number of rows
- `AVG()` gives the average value
- `MAX()` gives the maximum value
- `MIN()` gives the minimum value
- `SUM()` gives the sum of a range of values

There are lots of other functions that can let us manipulate text, understand more about our data, or perform more advanced operations

We can apply different functions to different types of data - for example, we can't use `AVG()` to calculate the average of a text type column



Practice

What is the highest price that the hospital pays for an item from the vendor with ID number 4?



Date Functions

- Use `YEAR()`, `MONTH()`, or `DAY()` to extract specific parts of dates
- Use `DATEDIFF(unit, date1, date2)` to calculate the difference between two dates
 - E.g. `DATEDIFF(year, 2013-06-13, 2020-06-13)` produces result 7
- Use `NOW()`, `GETDATE()`, or `SYSDATE()` to get the current date (differs based on RDBMS platform)



Practice

Generate a table containing `patient_id`, the patient's first and last names combined into one variable called `patient_name`, and the patient's allergies. The table should only include patients with an allergy to penicillin or morphine.



Today's 2nd Most Controversial Slide

JOINS vs. Subqueries



Joins

```
SELECT [columns to return]
```

```
FROM [left table]
```

```
[JOIN TYPE] [right table]
```

```
ON [left table].[field in left table to match] = [right table].[field  
in right table to match]
```



Left Join

		- - -	
		- - -	
		- - -	
		- - -	



Right Join

[illegible]



Inner Join

		- - -	
		- - -	
		- - -	
		- - -	
		- - -	



Right Join

SELECT

patients.patient_id,

patients.first_name,

patients.last_name,

admissions.room,

admissions.bed

FROM patients

RIGHT JOIN admissions ON patients.patient_id = admissions.patient_id;



Practice

All patients can access their medical documents on the hospital's website site. To access their documents, patients are given a temporary password. Show the `patient_id` and generate a password for each patient.

The password must be the following, in order:

1. `patient_id`
2. the numerical length of patient's `last_name`
3. year of patient's `birth_date`



Practice

We are looking for a specific patient. Pull all columns from the patient table for the patient who matches the following criteria:

- First_name contains an 'ian'
- Identifies their gender as 'F'
- Born in April, May, or November
- Their weight is between 80kg and 100kg
- Allergies field was unfilled in their patient record (NULL)
- They are from the city of Hamilton



Practice

Generate a table that shows the number of male patients from each province as `num_patients`. The province's full name (i.e. not the shortened `province_id`) must be shown in the table.





Today's Goal:

Be able to go into a relational database and pull out specific subsets of data using SQL.